

Per-interface DockQ, pre-trained vs 30-epoch fine-tuned

Each of 19 scored interfaces under both checkpoints, sorted by pre-trained score.

6 improved, 13 worsened; mean paired difference -0.015 (Wilcoxon $p = 0.16$).

Dashed line: acceptable threshold (DockQ 0.23). Two interfaces clear it under each model, but not the same two.

