

Training loss per epoch: re-run (Aug 18) vs reproduction run (Aug 17)

Rank-averaged mean diffusion loss, 30 epochs, 4 nodes / 16 GPUs, identical data and hyperparameters.

Dashed: linear fit. Re-run slope $-2.47\text{e-}4/\text{epoch}$ ($p = 0.021$); repro slope $-9.1\text{e-}5$ ($p = 0.17$); original record $-6.8\text{e-}5$ ($p = 0.49$).

Seeds are not fixed, so each run draws different noise; curves are separate samples of the same configuration.

