

Fixed-seed training loss: three repetitions nearly coincide

Rank-averaged mean diffusion loss per epoch, 4 nodes / 16 GPUs, torch.manual_seed(42 + rank).

Max pairwise epoch difference across the three seeded runs: 0.0019 (mean 0.0007); residual is nondeterministic-kernel rounding.

The unseeded Aug-18 run is shown for contrast: without a seed, curves are uncorrelated between runs.

