

Training loss over 30 epochs, 4 nodes / 16 GPUs

Rank-averaged mean diffusion loss per epoch. 60 gradient updates per rank, 1411 s.

Linear fit slope $-6.8\text{e-}5$ per epoch ($r^2 = 0.02$, $p = 0.49$): no significant decrease.

The training subset predates OpenDDE's 2021-09 data cutoff.

